

FPV-Tutorial - SS23

Materialien für Manuel's FPV-Tutorium im Sommersemester 2023

Manuel Lerchner

Zuletzt aktualisiert: 19. Juni 2023

FPV Tutorial - SS23

About

Materialien für Manuel's FPV-Tutorium im Sommersemester 2023.

- Die Materialien sind **privat** erstellt und können Fehler enthalten. Im Zweifelsfall haben immer die *offiziellen* Lehrunterlagen Vorrang.
- Alle Zusammenfassungen dieses Repositories können über manuellerchner.github.io/fpv-tutorial-SS23/summary.pdf heruntergeladen werden.
- Die Tutor-Slides sind unter manuellerchner.github.io/fpv-tutorial-SS23/slides.pdf verfügbar.

Found an error, or want to add something?

1. Fork this Repository
2. Commit and push your changes to **your** forked repository
3. Open a Pull Request to this repository
4. Wait until the changes are merged
5. A Github Action will automatically render the PDFs and deploy the static content to Github Pages

Contributors



Inhaltsverzeichnis

FPV Tutorial - SS23	1
About	1
Found an error, or want to add something?	1
Contributors	1
Week 1: Implications, Assertions and Stronges Postconditions	4
Implications	4
Definition of Implications	4
Truth Table	4
Examples	4
Assertions	5
Example for MiniJava	5
Strength of Assertions	6
Definition of Assertions-Strength	6
Special Assertions	6
Strongest Postconditions	6
Example	6
Week 2: Preconditions, Postconditions and Local Consistency	8
Weakest Preconditions	8
Example	8
Local Consistency	8
Example Local Consistency	9
Week 3: Loop Invariants	11
What is a loop invariant?	11
The Problem with Loops	11
Finding a Loop Invariant	11
Evaluating different Loop Invariants	12
Example Loop Invariants	12
Tips for Finding a Loop Invariant	12
Week 4: Termination Proofs	13
Why are Termination Proofs Necessary?	13
What is a Termination Proof?	13
How to do a Termination Proof?	13
Week 5: OCaml	15
Basic Syntax	15
Comments	15
Variables	15
Functions	15
Tuples	15
Lists	15
Records	15
If-then-else	15
Pattern Matching	15
Example Programs	16

Advanced Hello World	16
Debugging OCaml	16
Using the <code>#use</code> command in <code>utop</code>	16
Week 6: List Module	17
Lists in OCaml	17
Example Length	17
Binary Search Tree	17
Binary Search Tree in OCaml	17
Week 7: Advanced List Operations	19
List Reduce	19
List.fold_left	19
List.fold_right	19
Week 8: Lazy Lists	21
Lazy lists in OCaml	21
Partial Application	21
Week 9: Side Effects, Units and Exceptions	23
Side Effects	23
Example File IO	23
Exceptions	23

Week 1: Implications, Assertions and Strongest Postconditions

Implications

Implications are the key for understanding FPV. They show up in topics such as *Weakest Preconditions*, *Strongest Postconditions*, *Proof by Induction* / *Structural Induction*...

Definition of Implications

As you remember from the “Diskrete Strukturen” course, an implication is a statement of the form $A \implies B$. It is read as:

- “ A implies B ”
- “If A is true, then B is true”

It’s syntactic sugar for the following statement:

$$A \implies B \iff \neg A \vee B$$

This is a very important statement, because it can be used to simplify complex statements, if you can’t remember the specific rules for implications.

Truth Table

A	B	$A \implies B$
F	F	T
F	T	T
T	F	F
T	T	T

Examples

Example 1:

$$\begin{aligned} x = 1 &\implies x \geq 0 \\ \iff \neg(x = 1) \vee (x \geq 0) \\ \iff (x \neq 1) \vee (x \geq 0) \\ \iff \text{true} \end{aligned}$$

Example 2:

$$\begin{aligned}
& A \implies (B \implies A) \\
& \iff \neg A \vee (B \implies A) \\
& \iff \neg A \vee (\neg B \vee A) \\
& \iff \neg A \vee A \vee \neg B \\
& \iff \text{true} \vee \neg B \\
& \iff \text{true}
\end{aligned}$$

Assertions

Assertions are used to **annotate** specific points in a program and to **check** if a given expression is true at that point. If the expression is false, the program will terminate.

This is useful if you only want to allow certain values for a variable, because otherwise the program would not work as expected. They can also be used to prove the correctness of a program. Which is the main topic of this course.

Example for MiniJava



Abbildung 1: Flow Diagram

This corresponds to the following program:

```

void main() {
    var n = read(); //reads an arbitrary integer
    var i = 0;
    assert(A);

    var x = 0;
    assert(B)
    while (i < 10) {
        x = x + n;
        i = i + 1;
        assert(B);
    };

    write(x);
}

```

```

    assert(C);
}

```

The challenge is to find strong and precise assertions for the specific points in the program which allow us to prove the correctness of the program. In this case, we want to prove that the program **always** prints $n \cdot 10$ to the console. This corresponds to Assertion $C \iff x = n \cdot 10$.

Remember that whenever the program-flow reaches an assertion, the assertion must be true. Otherwise the program will terminate.

Strength of Assertions

Two assertions A and B can have different strengths. This happens for example if assertion A is more precise than assertion B .

For example, the assertion $A = 5$ is stronger than the assertion $B = 5 \vee 6$, because it is more specific. The assertion A only allows the value 5, while the assertion B allows the values 5 and 6.

This makes sense intuitively. But in order to use it in practice, we need to define what it means for an assertion to be stronger than another assertion.

Definition of Assertions-Strength

We say that an assertion A is stronger than an assertion B , if A implies B .

Using this definition, we can compare different assertions and determine if they are:

- **Equivalent:** $A \implies B$ and $B \implies A$
- **Ordered (eg. A is stronger):** $A \implies B$
- **Uncomparable:** $A \not\implies B$ and $B \not\implies A$

Special Assertions

Remember that *true* and *false* are also valid assertions. They are called **tautologies** and **contradictions** respectively.

How do they fit into the strength definition?

- **Tautologies:** $A \implies \text{true}$ for all A
 - This means that every assertion is stronger than *true* thereby making *true* the weakest assertion.
- **Contradictions:** $\text{false} \implies A$ for all A
 - This means that *false* is stronger than every assertion thereby making *false* the strongest assertion.

In practice those assertions show up in the following cases:

- **Tautologies:** If you have no information about the variables at a specific time in the program, you can use *true* as an assertion to express this.
- **Contradictions:** If you have a point that is **never** reached in the program, you can use *false* as an assertion to express this. The only way for the program to meet all assertions is to never reach such a point.

Strongest Postconditions

The strongest postcondition of a statement s and a precondition A is the strongest assertion B that holds after the statement s has been executed.

Example

Consider the following program:

```

void main() {
    var i=2;
    var x=6;
}

```

```

assert(x=3*i && i>=0);

i=i+1;

//state at this point:
//i = 3
//x = 6
//since the i in the assertions refers to the old value of i, before the statement i=i+1 was executed
//can we find a new assertion which explicitly computes the new value of x?
assert(C);
}

```

What is the strongest postcondition of the statement `i=i+1` and the precondition `x==3*i && i>=0`?
In other words what is the strongest assertion which we can insert in the second assertion?

This can be written as:

$$\mathbf{SP}[[i = i + 1]](x = 3 * i \wedge i \geq 0)$$

To compute the assertion after the statement `i=i+1` we basically need to **undo** the statement `i=i+1` because the original assertion referred to the old value of `i`, before it was updated.

Note: This only works for updates of variables. Other assignments might be a lot more complicated.

We first compute the **undo** of the statement `i=i+1`:

$$\mathbf{Undo}[[i = i + 1]] \equiv i = i - 1$$

Then we replace the variable `i` (which has already gotten updated) inside the assertion with the **undo-ed** statement:

$$\begin{aligned}
B &:= x = 3 * i \wedge i \geq 0 \\
&\longrightarrow x = 3 * (i - 1) \wedge (i - 1) \geq 0 \\
&\equiv x = 3(i - 1) \wedge i \geq 1 \\
&=: C
\end{aligned}$$

In total we have:

$$\begin{aligned}
C &:= \mathbf{SP}[[i = i + 1]](x = 3 * i \wedge i \geq 0) \\
&\equiv x = 3 * (i - 1) \wedge i \geq 1
\end{aligned}$$

Week 2: Preconditions, Postconditions and Local Consistency

Weakest Preconditions

Weakest Preconditions are used calculate the minimum requirements, which need to hold before an assignment, so that a given Assertion after the assignment holds.

Its written as:

$$\mathbf{WP}[[s]](e)$$

Where s is a statement and e is an assertion.

Example

Consider the following program:

```
void main() {  
    var r = 5;  
    assert(A);  
    var t = 3*r;  
    assert(t>=0);  
}
```

We want to find the minimal requirements which need to hold at `assert(A)` so that `assert(t>=0)` holds after `var t = 3*r;`.

We can calculate this using the following formula:

$$\begin{aligned} & \mathbf{WP}[[t = 3 * r]](t \geq 0) \\ & \equiv 3 * r \geq 0 \\ & \equiv r \geq 0 \quad \text{=: } A \end{aligned}$$

Now we know, that for the assertion `t>=0` to hold after `var t = 3*r;`, the assertion `r>=0` needs to hold before `var t = 3*r;`.

Local Consistency

Two assertions A and B are locally consistent, if A is **stronger** than the **weakest precondition** of B . This is written as:

$$A \implies \mathbf{WP}[[s]](B)$$

Note that it is not required that $A = \mathbf{WP}[[s]](\mathbf{B})$. Because a stronger assertion than required is also fine.

Local consistency is important: It mathematically proves that whenever the assertion A holds, then the assertion B holds after the statement s . This can be used to prove that a program actually computes what it is supposed to compute.

Example Local Consistency

Consider the following program:

```
void main() {
    var x = 30;
    assert(x>25); //A
    x=x+5;
    assert(x!=0); //B
}
```

$$\begin{aligned} A &\equiv x > 25 \\ B &\equiv x \neq 0 \\ s &\equiv x = x + 5 \end{aligned}$$

At the moment all the Assertions are arbitrary, and there is no guarantee that they actually hold during the execution of the program.

To prove them, we need to:

1. Show that all the assertions are locally consistent
2. We arrive at *true* at the start of the program

Local Consistency of A and B

We can calculate the weakest precondition of B and s as follows:

$$\begin{aligned} &\mathbf{WP}[[s]](\mathbf{B}) \\ &\equiv \mathbf{WP}[[x = x + 5]](x \neq 0) \\ &\equiv x + 5 \neq 0 \\ &\equiv x \neq -5 =: B' \end{aligned}$$

We can check the local consistency of A and B by checking if $A \implies B'$ holds.

This is the case, because:

$$\begin{aligned} A &\implies B' \\ &\equiv x > 25 \implies x \neq -5 \\ &\equiv \text{true} \end{aligned}$$

So we proved that A and B are locally consistent. This means that whenever A holds, then B holds after the statement s .

Weakest Precondition of A

If we compute the weakest precondition of A and $x=30$; we get:

$$\begin{aligned} &\mathbf{WP}[[x = 30]](x > 25) \\ &\equiv 30 > 25 \\ &\equiv \text{true} =: A' \end{aligned}$$

This is obviously also locally consistent, because $\text{true} \implies A' \equiv \text{true} \implies \text{true} \equiv \text{true}$.

Since we arrived at *true*, we know that the whole chain of assertions from the start to the end of the program holds and is locally consistent.

This means that we proved that when the assertion at the start (aka. *true*) holds, then the assertion *A* and consequently Assertion *B* holds.

In this case, we proved that in all instances of the program, the variable *x* cannot be 0 at the end.

Week 3: Loop Invariants

What is a loop invariant?

A loop invariant is an assertion that holds in each iteration of a loop. Finding such a loop invariant is needed to calculate weakest preconditions for programs with loops, because the **normal** way of finding the preconditions does not work for loops.

The Problem with Loops

Lets say you are trying to calculate the weakest precondition for the following program:

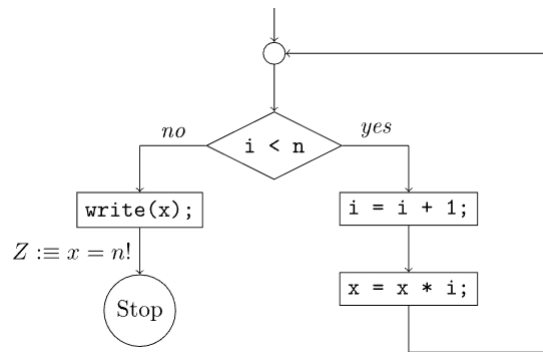


Abbildung 2: Program with loop

The normal way of finding the weakest precondition would be to start at the end of the program and work your way backwards.

$$X = \mathbf{WP}[\mathbf{write}(x)](x = n!) = x = n!$$

$$I = \mathbf{WP}[i < n](X, K) = (i < n \wedge K) \vee (i \geq n \wedge X)$$

$$K = \mathbf{WP}[i := i + 1](J) = J[(i + 1)/i]$$

$$J = \mathbf{WP}[x := x * i](I) = I[x * i/x]$$

So we came to a conclusion that in order for us to compute the weakest precondition I we need to calculate I, K and J . But J itself depends on I so we have a circular dependency and therefore we cannot calculate I directly.

Finding a Loop Invariant

Since we cannot directly compute loop invariants we need to find a way to come up with them indirectly.

We can do this by guessing a loop invariant I that is **strong enough** to prove the correctness of the program. We do this by checking that our assertions (which are constructed using our guessed loop invariant) are locally consistent.

If we have shown local consistency we just need to check if the starting point is annotated with *true*. Then we have successfully proven the correctness of the assertion at the end of the program.

Evaluating different Loop Invariants

For the program above a suitable loop invariant would be: $I := x = i! \wedge 0 < i \leq n$. But how do we come up with this loop invariant?

For this we look at some other loop invariants and evaluate them:

Example Loop Invariants

1. $I := x \geq 0$:
 - This loop invariant is **not strong enough** to prove the correctness of the program. Since it it fails the local consistency check. ($I \not\Rightarrow \mathbf{WP}[i < n](\mathbf{X}, \mathbf{K})$)
 - It was obvious that this loop invariant fails, because it does not contain any precise information about the value of x , which is needed to prove $x = n!$.
2. $I := i = 0 \wedge x = 1 \wedge n = 0$:
 - This loop invariant is way to strong, and is overall a bad choice because it would fail for any $n \neq 0$.
3. $I := x = i! \wedge 0 < i \leq n$:
 - This loop invariant is **strong enough** to prove the correctness of the program. Since it passes the local consistency check. ($I \Rightarrow \mathbf{WP}[i < n](\mathbf{X}, \mathbf{K})$)
 - Using this loop invariant we can prove that *true* holds at the start of the program, which means that the program and its assertions is correct.
 - Why does this loop invariant work?
 - It encapsulates all “relevant” information about the variables which change in the loop (x and i).
 - Combined with the false-branch of the if-statement it follows that $i \leq n \wedge i \geq n \Rightarrow i = n$. Which is exactly what is needed to prove $x = n!$ after we exit the loop.
 - It is weak enough to not disturb the proving of *true* at the start of the program.

Tips for Finding a Loop Invariant

There exist some old videos from the lecture “EIDI2” from the year 2017 that explain how to find loop invariants. The video is in german and is not relevant for this years course, but it still contains some useful tips for finding loop invariants.

- https://ttt.in.tum.de/recordings/Info2_2017_11_24-1/Info2_2017_11_24-1.mp4 [Nico Hartmann 2017]

Week 4: Termination Proofs

Why are Termination Proofs Necessary?

Every program containing a loop is potentially dangerous. Under the right circumstances, a loop can run forever, causing the program to hang. This is called an **infinite loop**.

In general, programs which don't eventually halt are of little use and can be considered as faulty.

For example, the following program contains an infinite loop and will never print "Finished":

```
let i = 17;
let j = 5;
while (i > j){
  i += 2;
  j += 1;
}
console.log("Finished");
```

On the other hand, the following program will always halt:

```
let i = 17;
let j = 5;
while (i > j){
  i += 1;
  j += 2;
}
console.log("Finished");
```

But how can we be sure that a program will always halt? In some cases it is not so obvious as in the examples above.

This is where **termination proofs** come in.

What is a Termination Proof?

In an assertion proof, you generally try to prove that a certain variable only takes on positive values inside a loop. Furthermore, you try to prove that the variable is decreased by at least one in each iteration of the loop.

This means that the variable will eventually reach zero (or less) and the cannot be entered again. Because this would violate the Assertion we defined.

But just coming up with arbitrary assertions and then claiming that they prove termination is not enough. We also need to show that those assertions are **locally consistent**.

How to do a Termination Proof?

Before we can perform a termination proof, it is necessary to understand what the loop actually does.

In the second example above our intuition tells us that the loop will eventually terminate. Because the variable j is increased by two in each iteration, while i is only increased by one. This means that j will eventually overtake i and the loop will terminate.

With this understanding we can define an auxiliary variable r which represents this intuition.

Since we only want to prove that the loop terminates, we just need to prove *true* at the end of the program.

In general, we need to insert the following assertions / statements



Abbildung 3: Flowchart with auxiliary variable

Notice that we need both these assertions to prove termination:

- $r > 0$ at the beginning of the loop
- $r > r_e$ at the end of the loop, right before $r = r_e$

Now the task is to show the following, we have proven that the loop terminates:

- local consistency of all assertions
- arrived at *true* at the start of the program
- The special assertions $r \geq 0$ and $r > r_e$ are also locally consistent

Week 5: OCaml

Basic Syntax

Comments

```
(* This is a comment *)
```

Variables

```
let x = 1;;
```

Functions

```
let f x = x + 1;;  
let succ: int -> int = fun x -> x + 1;;
```

Tuples

```
let x = (1, true);;
```

Lists

```
let x = [1; 2; 3];;  
let y = 1 :: 2 :: 3 :: [];;
```

Records

```
type person = { name: string; age: int };;  
let x = { name = "John"; age = 42 };;  
let name = x.name;;
```

If-then-else

```
let x = if 3 < 6 then "1" else "2";;
```

Pattern Matching

```
let x = match 3 with  
| 1 -> "1"  
| 2 -> "2"  
| _ -> "else";;  
  
let y = match "up" with  
| "up" -> (0, 1)  
| "down" -> (0, -1)  
| "left" -> (-1, 0)  
| "right" -> (1, 0)
```


Example Programs

Advanced Hello World

```
let welcome_string = "lmaCO ot emocleW ,!dlroW olleH"

let rec char_iterator sentence =
  if sentence = "" then ()
  else
    let head = String.get sentence 0 in
    let tail = String.sub sentence 1 (String.length sentence - 1) in
    char_iterator tail;
    print_char head

(* Call Function: Result will be visible in the console when loaded via utop *)
let _ = char_iterator welcome_string
```

Debugging OCaml

Using the `#use` command in `utop`

This method works best for single files containing the code you want to debug. The `#use` command works by just copy-pasting the entire content of the file into the `utop`-environment. It also shadows previous definitions of variables and functions.

1. Enter into the `root`-directory of the project and run `dune build` to initially build the project.

```
dune build
```

2. Open `utop` via `dune utop src`

```
dune utop src
```

In all projects you will encounter, `dune` is configured in a way that will allow this command to work.

This will open `utop` with the project's `src`-directory as the current working directory.

3. Reload the files
 - Instead of closing and reopening `utop` every time you change something in the files, you can use the `#use` command to load the files again.
 - But you need to be careful, because older definitions may still be around after reloading the file.

```
#use "src/main.ml";;
```

Week 6: List Module

Lists in OCaml

Lists in ocaml are basically linked lists. They are defined as follows:

```
let my_list = [1; 2; 3; 4; 5]
```

Since they are linked lists, you don't have direct access to the elements in the list. You can only access the head and the tail of the list. The head is the first element in the list and the tail is the rest of the list. The head is an element and the tail is a list. The tail can be an empty list.

```
let my_list = [1; 2; 3; 4; 5]
```

```
let head = List.hd my_list (* head = 1 *)
let tail = List.tl my_list (* tail = [2; 3; 4; 5] *)
```

Since Linked lists are recursive in nature, most of the functions that operate on lists are recursive. For example, the length function is defined as follows:

Example Length

```
let rec length l =
  match l with
  | [] -> 0
  | _ :: xs -> 1 + length xs
```

- This basically means that the length of list $l = \underbrace{[a_1, a_2, \dots, a_n]}_{\text{list}}$ can be recursively defined as
 - $\text{length}(\underbrace{[a_1, a_2, \dots, a_n]}_l) = 1 + \text{length}(\underbrace{[a_2, \dots, a_n]}_{\text{tail } l})$.
- Since we need a base case for the recursion, we define the length of an empty list to be 0.
 - $\text{length}([]) = 0$.

Many other functions on lists are defined recursively. For example, the reverse function is defined as follows:

Binary Search Tree

Binary Search Tree in OCaml

A binary search tree is a binary tree where the value of the left child is less than the value of the parent and the value of the right child is greater than the value of the parent. The following is an example of a binary search tree:

```
type tree =
  | Empty
  | Node of int * tree * tree

let my_tree =
  Node(6,
    Node(3,
```

```

        Node(1, Empty, Empty),
        Node(4, Empty, Empty)
    ),
    Node(8,
        Node(7, Empty, Empty),
        Node(9, Empty, Empty)
    )
)

```

The Tree type is defined as follows:

```

type tree =
  | Empty
  | Node of int * tree * tree

```

This means that a tree is either an empty node or a node with a left subtree, a value and a right subtree. The left and right subtrees are also trees. The value is an integer.

The syntax `Node of int * tree * tree` means that the constructor `Node` takes three arguments. The first argument is an integer and the second and third arguments are trees. The result of the constructor is a value of type `tree`.

Since Binary Search trees are again recursive in nature, most of the functions that operate on them are recursive. For example, the function `insert` is defined as follows:

Example Insert

```

let rec insert v t =
  match t with
  | Empty -> Node(v, Empty, Empty)
  | Node(y, l, r) ->
    if v < y then
      Node(y, insert v l, r)
    else if v > y then
      Node(y, l, insert v r)
    else
      t

```

It takes a value `v` and a tree `t` and returns a new tree with the value `v` inserted into the tree `t`. If the value `v` is already present in the tree `t`, then the same tree `t` is returned.

It works as follows:

- If the tree is empty, then the value `v` is inserted into the tree as the root node.
- If the tree is not empty:
 - If the value `v` is less than the value of the root node, then the value `v` is inserted into the left subtree.
 - If the value `v` is greater than the value of the root node, then the value `v` is inserted into the right subtree.
 - Else nothing is done and the same tree is returned.

As you can see from this example, it is not possible to change the value of a node in a tree. This is because the tree is *immutable*.

If you want to change the value of a node, you have to create a new tree with the new value.

This means, that if we need to insert a value in the right subtree of a node, we have to create a completely new tree with the new value inserted, and then replace the right subtree of the node with the new tree.

At first this seems unnecessary and inefficient, but in practise it is not that big of a deal. This is because the compilers of functional programs are in general very good at optimizing the code, since they can make more assumptions about the code due to the functional nature of the language. This means that the compiler can do a lot of optimizations that are not possible in imperative languages.

Week 7: Advanced List Operations

List Reduce

The `List` module contains a number of useful functions for working with lists. You can find the documentation for the `List` module [here](#).

Lets look at the `List.reduce` functions:

List.fold_left

Fold left is a function that takes a reducer function, an initial value, and a list. It then applies the reducer function to each element of the list, starting with the initial value, and then *eating* itself through the list from left to right.

It works just like the `Stream.reduce` function in Java

```
# List.fold_left (fun acc x -> acc + x) 0 [7; 8; -9; 10];;  
- : int = 16
```

If we split up the steps we get something like this:

list		7	8	-9	10
fold_left	acc	l[0]	l[1]	l[2]	l[3]
	0	7	8	-9	10
		7	8	-9	10
			15	-9	10
				6	10
					16

In each step we combine the accumulator with the next element in the list, and then use that as the accumulator for the next step.

It is called `fold_left` because it implicitly groups the elements of the list from left to right.

$$\begin{aligned} list &= [7, 8, -9, 10] \\ fold_left(f, 0, list) &= f(f(f(f(0, 7), 8), -9), 10) \end{aligned}$$

List.fold_right

`List.fold_right` is similar to `List.fold_left` except that it groups the elements of the list from right to left.

It has a different type signature than `List.fold_left`:

```
# List.fold_right (fun x acc -> acc + x) [7; 8; -9; 10] 0;;  
- : int = 16
```

list	7	8	-9	10	
fold_right	l[0]	l[1]	l[2]	l[3]	acc

list	7	8	-9	10	
	7	8	-9	10	0
	7	8	-9	10	
	7	8	1		
	7	9			
	16				

$$list = [7, 8, -9, 10]$$

$$fold_right(f, list, 0) = f(7, f(8, f(-9, f(10, 0))))$$

Week 8: Lazy Lists

A lazy list is a list that is not fully evaluated until it is needed. This is useful for representing infinite lists, or lists that are too large to fit in memory.

Examples of infinite lists include the list of all natural numbers, the list of all prime numbers, and the list of all Fibonacci numbers.

Lazy lists in OCaml

One way to implement lazy lists in OCaml is to use the `Lazy` module. But we are going to define our own lazy list type, which we will call `'a llist`.

```
type 'a llist = Cons of 'a * (unit -> 'a llist)
```

This type definition is a bit tricky. It says that a lazy list consists of a value of type `'a` and a function that returns another lazy list. The function is called a *thunk*. It is a function that takes no arguments and returns a lazy list. The thunk is used to delay the evaluation of the tail of the list.

Here is an example of a lazy list of natural numbers:

```
let rec from n = Cons (n, fun () -> from (n + 1))
```

It basically says that the lazy list `from n` consists of the value `n` and a thunk that returns the lazy list `from (n + 1)`.

Here is an example of a lazy list of Fibonacci numbers:

```
let rec fib_step a b = Cons (a, fun () -> fib_step b (a + b))
```

The parameters `a` and `b` are the two previous Fibonacci numbers. The lazy list `fib_step a b` consists of the value `a` and a thunk that returns the lazy list `fib_step b (a + b)`. Of course, the first two Fibonacci numbers are 0 and 1, so we can define the lazy list of all Fibonacci numbers as follows:

```
let fib = fib_step 0 1
```

This list is infinite, but since its generation is lazy, we can still make use of it.

Partial Application

Functions in OCaml are curried by default. This means that a function that takes two arguments is actually a function that takes one argument and returns a function that takes another argument. For example, the function `(+)` is defined as follows:

```
let (+) x y = x + y
```

This is equivalent to the following definition:

```
let (+) = fun x -> fun y -> x + y
```

The function `(+)` takes one argument `x` and returns a function that takes another argument `y` and returns the sum of `x` and `y`.

We can partially apply a function by passing it only some of its arguments. For example, we can define a function `inc` that adds 1 to its argument as follows:

```
let inc = (+) 1
```

```
let x = inc 5  
# val x : int = 6
```

The List module defines a function `map` that takes a function `f` and a list `xs` and returns a list of the results of applying `f` to the elements of `xs`. We can partially apply `map` to define a function `double` that doubles all the elements of a list:

```
let double = List.map (( * ) 2)
```

```
let xs = double [1; 2; 3]  
# val xs : int list = [2; 4; 6]
```

Also the `inc` function defined above can be used with `map`:

```
let xs = List.map inc [1; 2; 3]  
# val xs : int list = [2; 3; 4]
```

Which is equivalent to:

```
let xs = List.map ((+) 1) [1; 2; 3]  
# val xs : int list = [2; 3; 4]
```

Week 9: Side Effects, Units and Exceptions

Side Effects

In functional programming, side effects are a way of interacting with the outside world. For example, printing to the screen, reading from a file, or reading from the keyboard are all side effects. They are called side effects because they happen outside of the control of the program. In other words, the program cannot control when the user types something on the keyboard, or when the user clicks the mouse, or when data arrives from the network.

This also means that repeated calls to a function with side effects may produce different results. For example, if a function reads from the keyboard, then each time it is called, it may read a different value. This is in contrast to a function that does not have side effects, which will always return the same value when called with the same arguments. Those so called **pure** functions are the ones we have been writing so far.

Side effects are not necessarily bad. In fact, they are necessary for a program to be useful. However, they do make it harder to reason about the program. For example, if a function has side effects, then it is not enough to look at the function to understand what it does. You also need to know what side effects it has. This is in contrast to a pure function, which you can understand just by looking at the function itself.

Example File IO

Let's look at an example of a function with side effects. The following function reads the first line of a file and returns it as a string.

```
let read_file (filename : string) : string =  
  let ic = open_in filename in  
  let row = input_line ic in  
  close_in ic;  
  row
```

This function has side effects because it reads from a file. It is not a pure function because it does not always return the same value when called with the same arguments.

Writing to a file works similarly. The following function writes a string to a file.

```
let write_file (filename : string) (row : string) : unit =  
  let oc = open_out filename in  
  output_string oc row;  
  close_out oc
```

Exceptions

Exceptions are a way of handling errors. They are called exceptions because they are exceptional. They may occur in the normal course of a program, but they are not expected to occur. For example, if a function reads from a file, then it may fail if the file does not exist. This is an exceptional case because the file is expected to exist.

But sometimes we want exceptions to occur. For example if we read to the end of a file, then we want to know that we have reached the end of the file. Ocaml has a special exception for this case called `End_of_file`. We can use this exception to detect when we have reached the end of a file.

```
let rec read_file (filename : string) : string list =
  let ic = open_in filename in
  let rec read_loop () =
    try
      let row = input_line ic in
      row :: read_loop ()
    with End_of_file ->
      []
  in
  let result = read_loop () in
  close_in ic;
  result
```

The function `read_file` reads all the lines of a file and returns them as a list of strings.